

Introduction to Relational Databases

- Bachelor CS, Lille 1 University
- Nov. 16, 2011 (lecture 11/12)
- Today's lecturer: C. Kuttler
- Topic: Introduction to Database Tuning

1

Overview of today's lecture

- Introduction
- Query optimization
 - An example
- Tuning
 - Workload
 - Access structures
 - Operations (queries & modifications)
 - Modification of logical schema
 - Parameters of Architecture

2

Introduction

- Database design process
 - “logic design”: from the conceptual schema, we derive:
 - a logical schema
 - the necessary external schemas
 - “physical design”: starting from the logical schema, we derive:
 - a physical schema (access structures)
 - an optimized logical schema

3

Introduction

- Logic design
 - goal: design the database such that it avoids anomalies (normalization)
 - systematic algorithm exists for normalization
- Physical design
 - goal: all operations on the database are efficient
 - difficult to address systematically

4

Introduction

- Goal of physical design
 - is the performance
 - intervention on parameters that influence it
- Parameters with impact on the performance
 - organization of the files and access structures
 - logical schema
 - operations (queries and transactions)
 - parameters of the architecture (buffer, disks, etc)
- These aspects are difficult to plan

5

Introduction

- Typical approach to performance tuning
 - begin with the standard schema and the standard access structures
 - collect information on the actual use of the database, and evaluate the performance
 - based on these statistics, optimize the database's parameters (“tuning”)
 - repeat this activity is periodically

6

Introduction

- In today's class
 - Overview of physical design and tuning
 - We will discuss the main parameters
 - We will describe guidelines
 - But we can't be exhaustive
- Starting point
 - Evaluation and optimization of a relational DBMS

7

Query optimization

- How a query is evaluated
 - The query is either interactively sent to the DBMS, or sent by an application
 - The DBMS analyzes the SQL code syntactically
 - The DBMS checks the access permission
 - The DBMS performs the query optimization

8

Query optimization

- Optimization process
 - Choice of an efficient strategy for the query's evaluation
- Execution plan for a query
 - fix the order in which the necessary algebraic operators are applied
 - fix the strategy to compute the result of each operator through the available access structures (data structures)

9

Query optimization

In order to perform the optimization

- Alternative execution plans are evaluated
- The optimizer uses statistics on the database's content
 - dimensions of the tables, dimensions of the records, dimensions of the indices, selectivity, etc
- The cost of each execution plan is estimated, based on these statistics
 - number of block accesses on the hard disk

10

Example: university database

```
CREATE TABLE teacher (  
  t_id char(4) PRIMARY KEY,  
  lastName varchar(20) NOT NULL,  
  firstName varchar(20) NOT NULL,  
  qualification char(15),  
  school char(10) );
```

```
CREATE TABLE student (  
  s_id integer PRIMARY KEY,  
  lastName varchar(20) NOT NULL,  
  firstName varchar(20) NOT NULL,  
  program char(20),  
  year integer,  
  advisor char(4) REFERENCES teacher(t_id)  
);
```

11

Our example

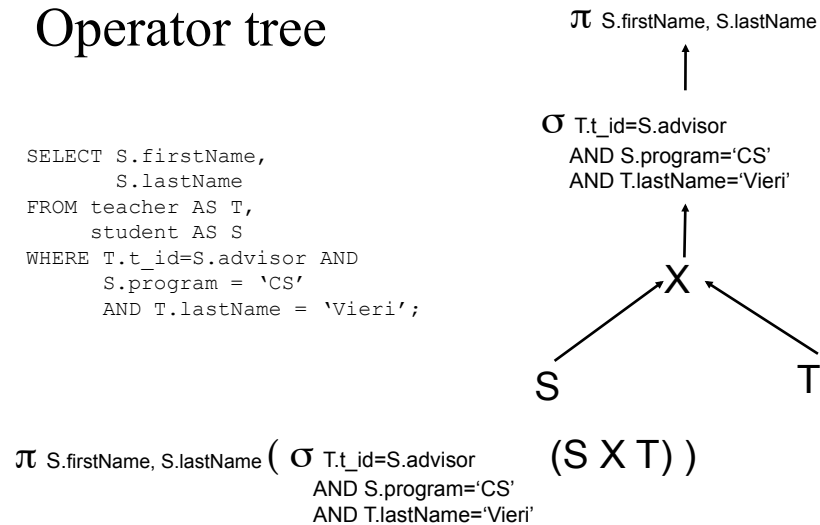
- Query: *“First and last names of Prof Vieri's students in Computer Science”*

```
SELECT student.firstName,  
       student.lastName  
FROM teacher T, student S  
WHERE T.t_id=S.advisor AND  
       S.program = 'CS' AND  
       T.lastName = 'Vieri';
```

12

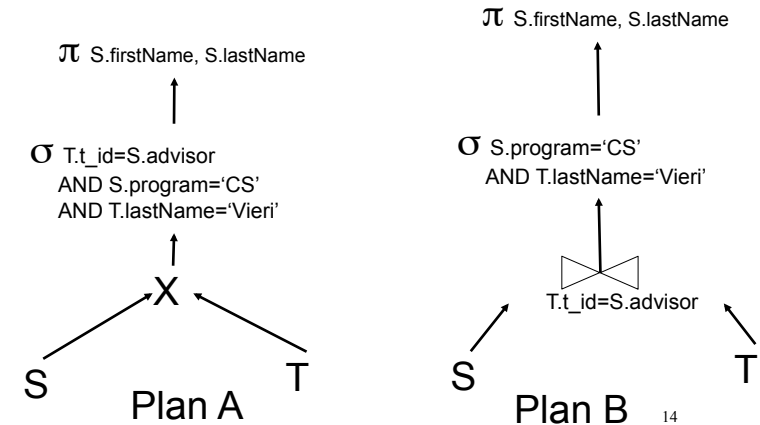
Operator tree

```
SELECT S.firstName,
       S.lastName
FROM teacher AS T,
     student AS S
WHERE T.t_id=S.advisor AND
      S.program = 'CS'
      AND T.lastName = 'Vieri';
```



13

Same example – other trees



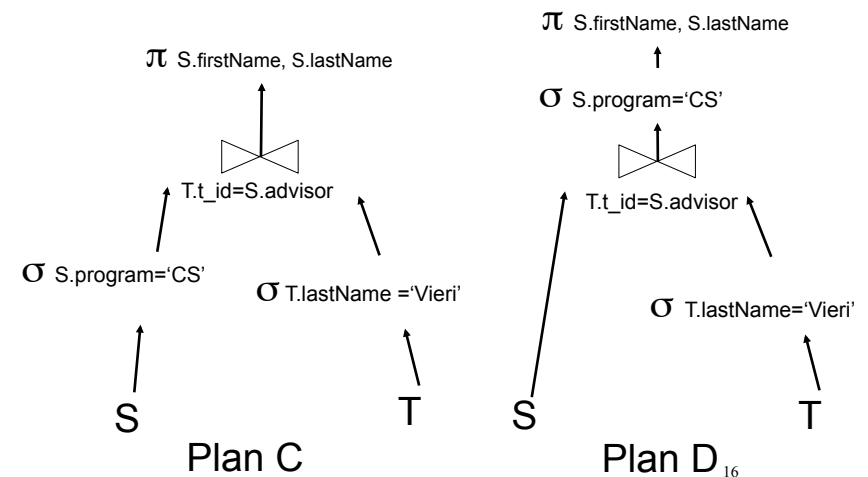
14

Which is better, A or B?

- Generally, the optimizer would prefer the plan B to the plan A
 - Cartesian products are expensive
- But in other cases
 - the operator tree alone isn't enough to know if one strategy is better than the other
 - the available access structures must be considered

15

Other execution plans



16

Which is better, C or D?

- Generally
 - plan C is better than plan D
 - but under certain conditions, plan D is better
- It depends on the strategy to evaluate the operators
 - in particular, the file organization
 - the data structures for access (indices)

17

Execution of algebraic operators

- Three main techniques
- Linear scan of the file
 - Inefficient, only applied to small files
- Access through index
 - Assumes the presence of indices (up to date)
- Temporary grouping
 - Creation of additional structures to group the tuples (example: order, hash table in main memory)

18

Execution of a selection: the data structure matters!

- Unordered file, no relevant index available
 - linear scan of the file
- File ordered by the attribute, no index
 - binary search in the file
- Index B+-tree on the attribute
 - B+-tree: data in external nodes, keys in internal nodes. Internal nodes have up to b children.
 - search in the index with complexity $O(\log_b n)$
- Hash index on the attribute
 - direct access through hash function, ideally

19

Execution of a join

- Basic strategy, without index
 - Embedded cycles
 - Pretty inefficient...
- Example: $S \text{ JOIN } T \text{ ON } T.t_id = S.advisor$
 - for each tuple of S
 - for each tuple of T
 - if $T.t_id = S.advisor$ then output the resulting tuple

20

Execution of a join

- Embedded cycles with indices
 - Exploit an index on one of the join attributes
- Example: **S join T on S.advisor = T.t_id**
 - index on T.t_id
 - For each tuple of S
 - for each tuple of T such that S.advisor=T.t_id
 - include the tuple in result

index based access
21

Execution of a join

- “Sort-Merge” Join
 - idea: if both tables are ordered on the join attribute, the join is linear
- Strategy
 - Create an ordered copy of the table
 - Generate the result by a linear scan
 - Particularly efficient if table is already ordered

22

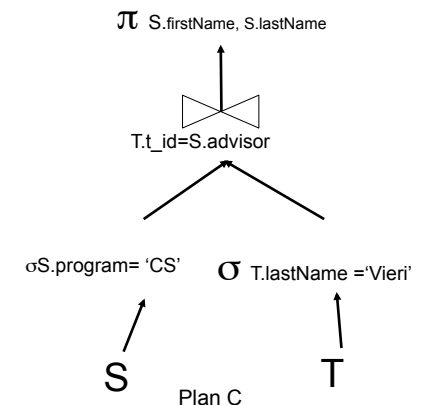
Execution of a join

- Hash Join
 - idea: hash on the join attribute for both tables
- Strategy
 - in main memory, build hash tables for both tables on the join attribute
 - Scan a table, and for each value, use the hash function to localize the corresponding bucket of tuples

23

Plan C

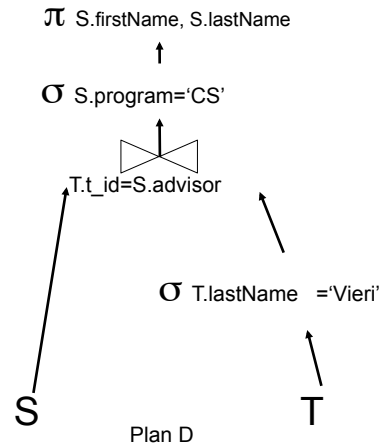
- Assume that:
 - no relevant indices
- Selection
 - linear scan
- Join
 - sort-merge
- Complete plan



24

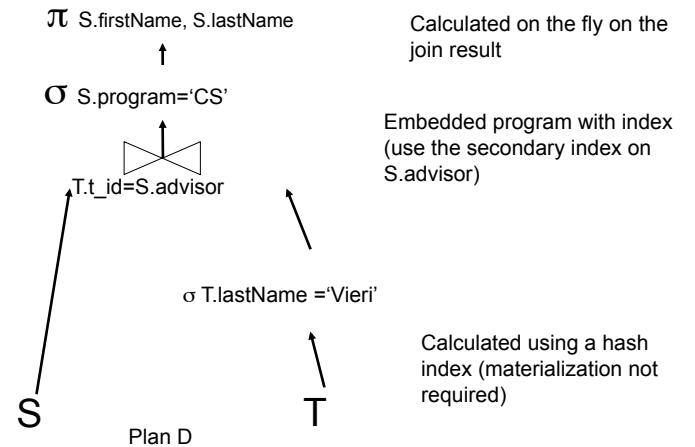
Plan D

- Suppose:
 - Hash index on T.lastName
 - Hash index on S.advisor
- Selection
 - hash
- Join
 - Cycles with index



25


Complete plan D



26

Plan C, other hypothesis

- Assume:
 - Hash index on T.lastName
 - Secondary index on S.program
- Not very useful for S
 - because the number of different programs is low
 - low selectivity of the index



27

Optimization

- Top-end DBMS allow to consult the selected execution plan
- EXPLAIN command
 - Typical syntax: EXPLAIN <select>
 - Illustrates the execution plan and estimates its cost, by the optimizer
 - Available in PostgreSQL, MySQL, ...

28

Tuning

- Typical scenario
 - After an initial phase, the actual use of a database is analyzed
 - The performance isn't good
 - An intervention is needed to improve the performance, by improving the parameters
- Starting point
 - Workload: usage statistics for the database

29

Workload

- Optimization isn't possible for all possible queries
- Only the most frequent and relevant operations are considered
- Workload
 - List of queries
 - List of updates
 - For each, expected performance (for example: <2s, or number of transactions per minute)

30

Tuning activities

1. Choice of access structures
 - Organization of files, indices, clustering
2. Interventions on operations
 - Re-structuring, isolation level
3. Interventions on the logical schema
 - Partitioning, aggregation, denormalization
4. Interventions on parameters of the architecture
 - buffer, disks, RAM

31

Access structures

- Main kind of intervention
 - Addition of indices
- Be careful
 - indices improve the performance
 - But slow down updates
 - Need disk space
 - A compromise is needed

32

Access structures

- Extreme case: read-only database
- Example
 - Archive of the immigrants in the USA between 1800 and 1900, with millions of tuples
 - Search by lastName, firstName and arrival year
- Many indices can be used
 - lastName, firstName, year, lastName and year, lastName and firstName

33

Access structures

- Guideline 1
 - It is only worth introducing an index, if it allows to improve the performance of more than one query of the workload
- Warning:
 - The optimizer doesn't always manage to use an index
 - example: `select * from Employees where AnnualSalary/12 > 3000`
 - Check the execution plan before and after

34

Access structures

- Guideline 2
 - The attributes to intervene on are those that appear in joins and selections
 - For equality checks (ex: `income=5000`), hash indices are preferable
 - For conditions on intervals (ex: `income > 5000` and `income < 10000`) B+-tree are preferable

35

Access structures

- Guideline 3
 - It is only worth introducing an index if the number of different values of an attribute is sufficiently high
- Example:
 - `select * from employees where income = 10000`
 - The index on income might not be useful, if many employees have the same income

36

Access structures

- Guideline 4
 - Pay attention to bottlenecks
- Example 1:
 - Unordered relation with frequent insertions
 - The last block is a bottleneck
- Example 2:
 - Schema modifications (lock in writing into the catalogue)

37

Intervention on operations

- Two forms of interventions on operations
 - Restructuring queries
 - rewrite the query in a smarter way, such that its execution can be optimized
 - after some experience on the job, you should come back to this topic!
 - Choice of isolation level for transactions

38

Intervention on operations

- Example: restructuring a query
 - `select * from Employee`
 `where yearlyIncome / 12 > 3000`
 - `select * from Employee`
 `where yearlyIncome > 3000 * 12`
- Other forms of restructuring
 - limit the level of embedding (difficult to optimize)
 - ...

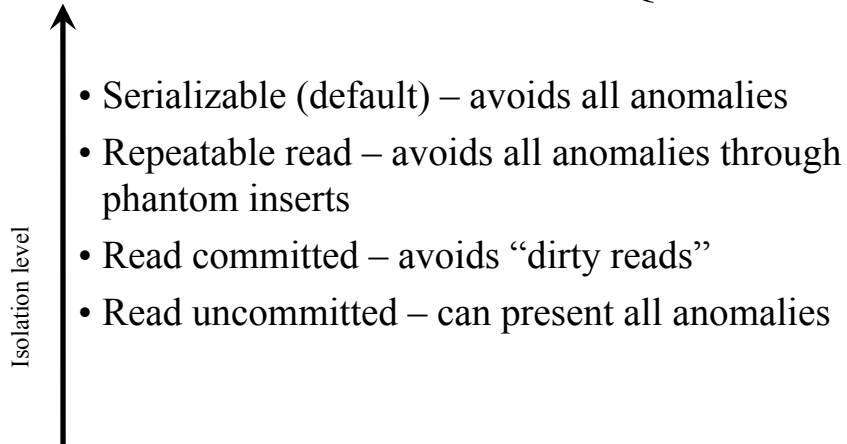
39

Intervention on operations

- Isolation level for transactions
 - The usual level is `SERIALIZABLE`
 - Often, `READ COMMITTED` is adequate
 - Generally, it is useful to separate interactive queries and updates

40

Isolation levels in SQL



41

Modification of logical schema

- A normalized schema is not necessarily more efficient
 - normalization: formal method to reach “best” possible schema, that avoids most anomalies
- Four main forms of interventions
 - Partitionning tables
 - Aggregation of tables
 - Denormalization of tables
 - Adding redundant information

42

Modification of logical schema

- Be careful
 - When the logical schema is modified, some applications may not work correctly
- Two possible solutions
 - The modifications of the logical schema are decided very early (immediately after logical modeling)
 - Or, if possible, one creates an external schema equal to the old logical schema

43

Modification of logical schema

- Partitioning tables
 - Tables with many attributes can be split in two smaller tables
- Example: the table “student”
 - Primary key (s_id)
 - Personal data (firstName, lastName, s_id, social security number, address, family income etc.)
 - Academic attributes (program, year, advisor, internship, company etc.)

44

Modification of logical schema

- Can break the table into two tables
 - StudentPersonalData: s_id and all personal attributes
 - StudentUniversityData: s_id and all academic attributes
- Useful when
 - one rarely needs to access both kind of data
 - in those cases, a join will be needed

45

Modification of logical schema

- Warning
 - such restructuring must be done very early
 - views don't help
 - defining the view “student” as the join of the two tables, wouldn't improve the performance

46

Modification of logical schema

- Aggregation of tables: combining two tables into one
 - may avoid joins
- Example: student and internship
 - External key “s_id” of internship
 - If the internship data are frequently accessed, it makes sense to combine both tables into one
 - This avoids joins
 - The number of null values increases
 - Can define two views to present the schema

47

Modification of logical schema

- Denormalization of tables
 - Normalization avoids anomalies, but often enforces too many joins
- Example: teachers and numbers
 - Numbers(number, teacher FK)
 - If we often need to print the list of names and ids, then we could add the teacher's firstName to the table Number
 - However, this increases the complexity of updates

48

Modification of logical schema

- In this case
 - We generate (slight) update anomalies
 - example: each time that a teacher's lastName changes, need to intervene as well on teacher as on Numbers
 - To avoid to create inconsistent instances of the database, one needs to use transactions

49

Modification of logical schema

- Adding redundant information
 - Can sometimes avoid complex queries
- Example: number of exams that have been taken by a student
 - Can be derived by aggregation of the join between students and exams
 - Can be explicitly stored as an attribute of student
 - Forces to use transactions

50

Parameters of architecture

- Buffer
 - Increasing the buffer increases the "hit ratio"
 - Makes sense up to a certain limit
- Disks
 - Placing the files on several disks improves the performance
 - example: disk for the log (the log is a typical bottleneck)

51

Summary

- Introduction
- Query optimization
 - An example
- Tuning
 - Application workload
 - Access structures
 - Modification of logical schema
 - Architectural parameters

52